

Real-Time Fraud Detection on Credit Cards Transactions

Prasanna Surianarayanan, Dipika Rajesh

ps3703@nyu.edu, dr2898@nyu.edu


Premise

- Fraud detection applies to a variety of sectors:
 - Banking and Finance
 - Insurance
 - Government agencies and law enforcement
- 2018 PwC Survey : fraud experienced by 49% of the 7,200 companies surveyed
- Fraud Detection : A billion-dollar business

Need for Credit Card Fraud Detection

- Electronic commerce has rapidly increased over the years
- Increasing confidence of customers in electronic payments
- Customers increasingly vulnerable to new attacks
- Frauds need to be detected in (near) real time
- Scalable learning techniques needed to ingest and analyse massive amounts of streaming data.

Credit Card Fraud Reports in the United States




Leveraging Big Data for the Use Case

- Credit card transactions captured over time based on the cardholder's spending behaviour
- Events stored as transaction history for detection of fraud in the future
- Spark Structured Streaming and Apache Kafka work pretty well in capturing such incidents real time
- Ecosystem capable of handling vast data with good throughput and low latency
- SparkML used to develop the prediction model

Technologies Used

- Apache Kafka for message ingestion
- Spark Structured Streaming to process datastreams from Kafka
- Spark ML to detect fraudulent transactions

Proposed Architecture



Ecosystem Overview

- **Apache Kafka :**
 - Transactions ingested to Kafka
 - Used for delivering streams of card transactions
 - Card transactions are published to a topic
 - Consumers read messages from subscribed topics

Ecosystem Overview

- **Spark Structured Streaming :**

- Reads stream of data from the Kafka topic
- Creates dataframes in streaming mode
- Treats arriving data as an unbounded input table
 - New items in streams treated like rows, appended to the input table
- Queries applied to streaming dataframes just like static dataframes
- Triggers used to control result updates

Ecosystem Overview

- **Spark ML :**
 - Used to create Models to detect fraudulent transactions
 - Historical transactions along with the class labels are used to train the model
 - Prediction is performed on the incoming data from the Kafka topic
 - Predicted results written to an output sink


Dataset Used

- Source of labelled data : Kaggle - <https://www.kaggle.com/mlg-ulb/creditcardfraud>
- Total features : 31
- 28 features transformed using PCA for confidentiality purposes
- 3 Non-anonymized features : Time, Amount, Class Label
- Dataset was selected due to non-synthetic, high-dimensional data.

Data Exploration!


Dimensionality Reduction using t-SNE

T-distributed Stochastic Neighbor Embedding (t-SNE) nonlinear dimensionality reduction technique well-suited for embedding high-dimensional data for visualization in a low-dimensional space of two or three dimensions




Exploration of the Dataset

The vast majority of transactions are relatively small and only a tiny fraction of transactions comes even close to the maximum - needs to be scaled!

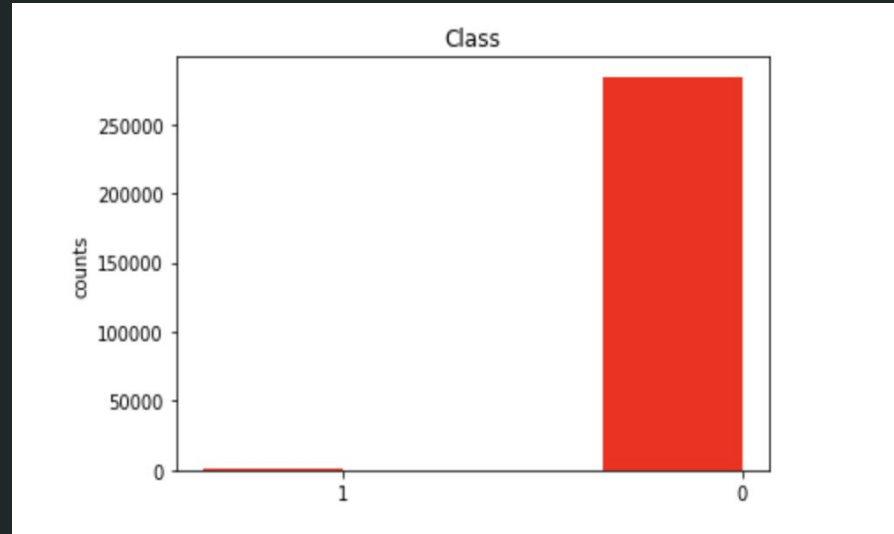


- The time is recorded in the number of seconds since the first transaction in the data set.
- Can conclude that this data set includes all transactions recorded over the course of two days.



Class Distribution?

Most transactions are non-fraudulent. In fact, 99.83% of the transactions in this data set were not fraudulent while only 0.17% were fraudulent.




Data Preprocessing!

- Imbalanced datasets can create biased machine learning models
 - Oversampling methods (SMOTE) were applied on the dataset along with feature scaling.
-

After applying SMOTE,

the oversampling created
balanced classes!




Choosing and Fitting the Machine Learning model

Model Choice - Random Forest Classifier

Advantages of Random Forest Classifiers:

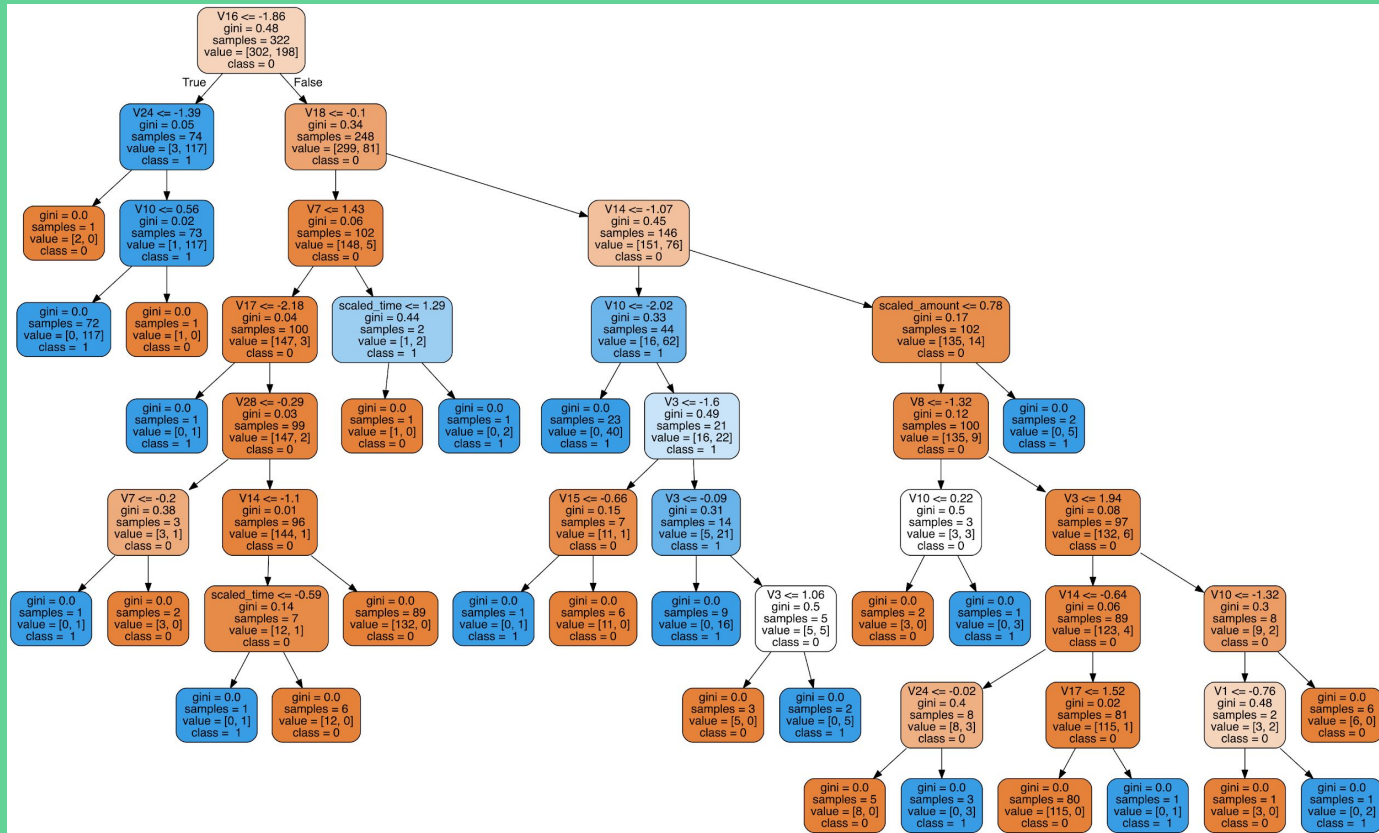
- Generalizes patterns well
- Robust to different input types
- Robust to outliers
- Highly scalable



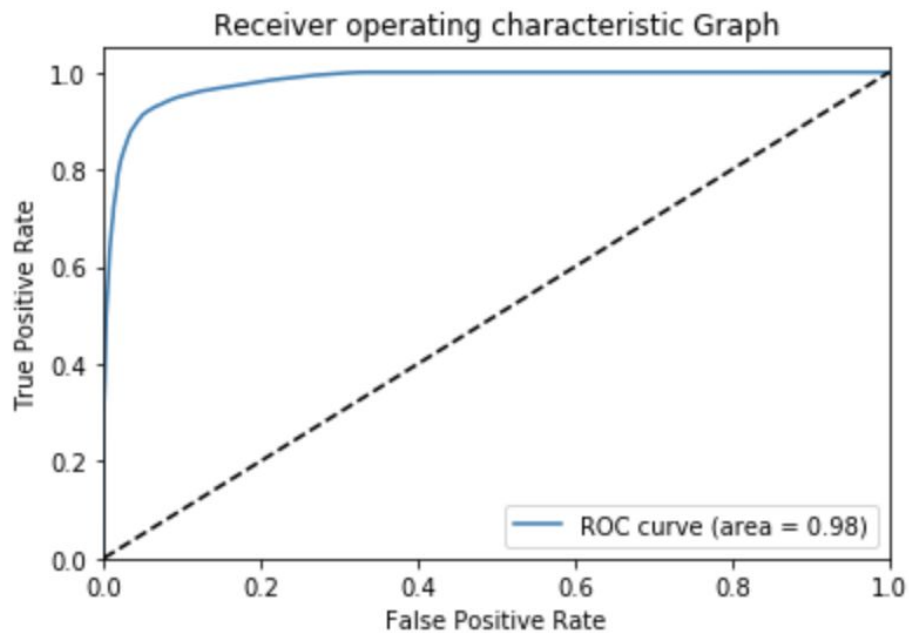
Using RandomForestClassifier
from SparkML library, trained on
NYU DUMBO HPC with 500,000
data instances!

```
algo = RandomForestClassifier(featuresCol='features', labelCol='Class')
start_time = time()
model = algo.fit(trainingData)
end_time = time()
predictions = model.transform(testData)
```

Visualizing the Random Forest model using Eureka Trees library




Results! - Area Under the Receiver-Operating Characteristic (AUROC) metric




AUC — ROC Curve

Evaluation of the random forest model - confusion matrix



Recall in the train set = 1.0



Recall in the test set = 0.8095

Never use accuracy score as an evaluation metric on imbalanced datasets - will be high and misleading.

Never test the model on the sampled (over or under) data.

Future Work

- Adding ElasticSearch and Kibana to the pipeline for visualization and trend analysis
 - Expanding model to include Time-Series Analysis of Customer Credit Card Transactions
 - Online Machine Learning with incoming real-time data
-

Thank you!